

15-381

Artificial Intelligence



Drew Bagnell and Emma Brunskill

Fall 2013


Multiple slides over course taken or
adapted from Dan Klein, and past
CMU 15-381 Instructors



Can Machines Think?

Alan Turing, “Computing Machinery and Intelligence,” *Mind*, 1950.

Mind-Body Dualism



The Astonishing Hypothesis

“You, your joys and your sorrows,
your memories and your ambitions,
your sense of personal identity and
free will, are in fact no more than the
behaviour of a vast assembly of
nerve cells and their associated
molecules”

--- Francis Crick, 1994

“You insist that there is something a machine cannot do. If you will tell me precisely what it is that a machine cannot do, then I can always make a machine which will do just that.”

--- John von Neumann, 1948

What is intelligence?


“...making a machine behave in ways that would be called intelligent if a human were so behaving.”

--- John McCarthy, 1955

“One might ... define thinking as consisting of those mental processes that we don't understand.”

--- Alan Turing

Imitation Game: Turing Test



Answer: X is a human and Y is a machine

Chinese Room

- Thought experiment proposed by Searle
- Suppose AI has produced a program that can pass the Turing Test in Chinese
- You have a handbook with its pseudocode
- You're in a closed room and receive Chinese characters through a slot
- You run the program's code manually and return the output
- Does this mean you *understand* Chinese?

Counterarguments

- Finding the mind: the whole system understands Chinese, the person is just a part of the system
- Redesigning the experiment: suppose the program simulates the actions of every neuron in the brain of a Chinese speaker

Counterarguments

- Doubts about the intuition
 - Brain performs 100 billion operations per second, so it would take the person millions of years to simulate a simple answer
 - Churchland's Luminous Room: suppose you are standing in a dark room and quickly moving a magnet up and down, then by Maxwell's theory of artificial luminance it will be luminous. However, this requires 450 trillion movements per second

DARTMOUTH CONFERENCE (1956)

Founding of AI

“The study is to proceed on the basis of the conjecture that every aspect of learning or any other feature of intelligence can in principle be so precisely described that a machine can be made to simulate it.”


--- John McCarthy

DARTMOUTH CONFERENCE (1956)

Founding of AI




John McCarthy



Marvin Minsky



Claude Shannon



Oliver Selfridge


and “two vaguely known persons from RAND and Carnegie Tech... a significant afterthought.”
--- Pamela McCorduck, “Machines Who Think”

DARTMOUTH CONFERENCE (1956)

Founding of AI




John McCarthy



Marvin Minsky



Claude Shannon



Oliver Selfridge




Alan Newell

Herbert Simon

**Act like
a Human**


**Act
Rationality**

- **Rational:** Optimal decision making, given current information, towards achieving a goal.



WHAT CAN AI DO NOW?

Play a decent game of table tennis?



WHAT CAN AI DO NOW?

- ✓ Play a decent game of table tennis?
- Drive safely along a curving mountain road?
- Drive safely along Forbes?
- Buy a week's worth of groceries on the web?
- Buy a week's worth of groceries at Giant Eagle?
- Discover and prove a new mathematical theorem?
- Converse successfully with another person for an hour?
- Fold laundry?
- Perform a complex surgical operation?
- Unload a dishwasher and put everything away?
- Translate spoken Chinese into spoken English in real time?
- Write an intentionally funny story?

WHAT CAN AI DO NOW?


- ✓ Play a decent game of table tennis?
- ✓ Drive safely along a curving mountain road?
- ✓ Drive safely along Forbes?
- ✓ Buy a week's worth of groceries on the web?
- ✗ Buy a week's worth of groceries at Giant Eagle?
- ? Discover and prove a new mathematical theorem?
- ✗ Converse successfully with another person for an hour?
- ✓ Fold laundry?
- ? Perform a complex surgical operation?
- ✓ Unload a dishwasher and put everything away?
- ✓ Translate spoken Chinese into spoken English in real time?
- ✗ Write an intentionally funny story?

DARPA Urban Challenge

- 96 km urban area course, to be completed < 6 hours, took place in 2007
- Tartan Racing (CMU+GM) claimed the \$2 million prize
- Challenge involves mission planning, motion planning, behavior generation, perception, world modeling
- <http://www.youtube.com/watch?v=IULI63ERek0>

Watson

- Watson defeated the two greatest-ever Jeopardy! champions
- Involves natural language processing, information retrieval, knowledge representation and reasoning, and machine learning
- http://www.youtube.com/watch?v=oUj9AzSE_9c




MORAVEC'S PARADOX

The main lesson of thirty-five
years of AI research is that
the hard problems are easy and
the easy problems are hard.


--- Steven Pinker

Evolution of Intelligence


- Human genes: 30 billion base pairs x 2 bits = 750 megabytes
- Human brain: 100 billion neurons x 1 bit = 12.5 gigabytes
(a certain underestimate)
 - Better (but still under-) estimate
 - 10^{14} synapses \rightarrow 12,500 gigabytes (**G. Orr**)
 - Internet? ~500 exabytes (IDC, 2009)




Biological Information Storage




Credit: Carl Sagan, *The Dragons of Eden*, 1977



Gene



Brain



Extrasomatic Knowledge



RAM



Hard drive



External Storage

Biology vs Engineering



“AI in the Postmodern Age”

--- Alexei Efros

- All interesting questions have been asked
- **Intelligence:**
 1. Remember all questions, the best answers
 2. Match current question
 3. Look up the best answer



A screenshot of a Google search interface showing search suggestions. The user has typed "Artif" into the search bar, and a dropdown menu lists the following suggestions:

- Artif
- artificial intelligence
- artificial life
- artificial insemination
- artifact
- artifice
- artificial grass
- artificial girl 3
- artificial selection
- artificial sweeteners
- artificial heart

At the bottom of the suggestions box are two buttons: "Google Search" and "I'm Feeling Lucky".



WolframAlphaTM computational...
knowledge engine

Enter what you want to calculate or know about:

What is the meaning of life



[≡ Examples](#) [Random](#)

Assuming "What is the meaning of life" is a quantity | Use as referring to English words instead

Input interpretation:

Answer to the Ultimate Question of Life, the Universe, and Everything

Result:

42

(according to Douglas Adams' humorous science-fiction novel *The Hitchhiker's Guide to the Galaxy*)

Computed by **Wolfram Mathematica**

Download as: [PDF](#) | [Live Mathematica](#)

Give us your feedback:



[About](#) | [Products](#) | [Mobile Apps](#) | [Business Solutions](#) | [For Developers](#) | [Resources & Tools](#)

[Blog](#) | [Forum](#) | [Participate](#) | [Contact](#) | [Connect](#)



WolframAlpha™ computational...
knowledge engine

log16(round((mass of the universe)/(protein of a large egg))) - 3



[≡ Examples](#) [» Random](#)

Assuming egg | Use [raw egg](#) instead

Assuming any type of egg | Use [egg](#), [fried](#) or [more](#) instead

Input interpretation:

$$\log_{16} \left(\text{Round} \left[\frac{\text{mass of the observable universe}}{\begin{array}{c} \text{egg} \\ \hline \text{amount} \end{array} \quad \begin{array}{c} 1 \text{ large egg} \\ \hline \text{protein} \end{array}} \right] \right) - 3$$

$\log_b(x)$ is the base b logarithm »

Result:

42

[Show details](#)

Computed by **Wolfram Mathematica**

[Source information](#) »

[Download as: PDF](#) | [Live Mathematica](#)

Give us your feedback:

[send](#)

Artificial Artificial Intelligence

- Amazon mechanical turk
- Crowdsourcing
- But there also exists AI to improve artificial AI...

Representation

- All knowledge needs representation
 - Genetic Information
 - Behavioral Information
 - Audio Information
 - Visual Information
 - Textual Information

A major part AI is representing the problem space to allow efficient optimization for the best solution

Technological Singularity

- Emergence of superhuman intelligence
- Key idea: self-improvement
- Source of name: analogy between inability to predict events after the development of a superintelligence, and the space-time singularity beyond the event horizon of a black hole
- Some predict: this century
- Others argue: never



Understanding & creating
intelligence
is the grand intellectual
challenge of the century

Rene Descartes
Philosopher...

Warren McCulloch
Neurophysician

Francis Crick
Molecular Biologist

Marvin Minsky
Cognitive Scientist

Aristotle
Philosopher...

John von Neumann
Mathematician...

Herbert Simon
Economist...

Claude Shannon
Information Theorist

Thomas Bayes
Reverend

Alan Turing
Computer Scientist

Walter Pitts
Logician

Alan Newell
Cognitive Psychologist

“AI still has openings
for a full-time Einstein”
--- Russell and Norvig

“If you invent a breakthrough in AI,
so machines can learn,
that is worth 10 Microsofts.”
--- Bill Gates

Our Expectations From You

AI COURSE TOPICS

- Making good decisions in deterministic worlds
 - Search
 - Constraint satisfaction
 - Optimization
- Handling Uncertainty
 - Bayes' nets
 - Machine learning
 - Decision making under uncertainty
- Beyond 1 mind: multi-agent AI
- Applications (robotics, vision...)

ARTIFICIAL INTELLIGENCE

Administration

- **Instructors:** Drew Bagnell and Emma Brunskill
- **TAs:** Jesse Dunietz, Subhodeep Moitra, Danny Zhu
- **Course Website:** : <https://www.cs.cmu.edu/~15381/>
- **Class Timing:** Tuesday and Thursday, noon-1:20pm
- **(Strongly) Recommended Book:** Artificial Intelligence: A Modern Approach (Russell and Norvig)
 - Will try to make copies available at library as well

ARTIFICIAL INTELLIGENCE

Administration

- **Reading:** Expected to read **in advance** of lecture
- **Question answering through Piazza**
- **Homework:** Can use 8 late days, at most 2 per homework
- **First homework will be due 9/10**
- **No laptops or cell phones**
 - We expect total participation in class
- **Extra credit:** Participation through
 - Clickers/question response
 - in-class question/answer

Background Calibration

- Rest of today: getting a sense of background
- We don't expect you to know all (or perhaps any) of these things
 - Very diverse background in class
 - **Only graded on participation**
 - Just gives us a better sense of everyone's background knowledge so we can consider recitations