

15-451 Assignment 07

Karan Sikka

ksikka@cmu.edu

Collaborated with Dave Cummings and Sandeep Rao.

Recitation: A

November 7, 2014

1: A Densely-Knit Community

(a) Consider the following assignment of values to x :

Let $x_i = 0$ if v_i not in S^*

Let $x_i = \frac{1}{|S^*|}$ if v_i is in S^*

Let $x_{ij} = \min(x_i, x_j)$

Note that the constraints are satisfied. All values are at least 0. The sum of the x_i is 1 since it's $|S^*| * \frac{1}{|S^*|}$.

$$\text{OPT}_{LP} \geq \sum_{i,j \in E} x_{ij} = \text{edges}(S^*) * \frac{1}{|S^*|} = \text{OPT}$$

(b)

(c)

(d)

2: Large + Dense = Difficult

The problem is in NP because there exists a poly-time verifier as follows:

Given a (G, K, D) and a potential solution S of size K , compute the density which is

$$\frac{|\text{edges}(S)|}{|S|}$$

Note: Computing edges of S can be done in $|S|^2$ time.

The problem is NP-hard because we can show a Karp reduction from K-Clique to it.

Given G, K we want to output YES if there exists a K-clique in G .

Observe that the number of edges in a K-clique is $\frac{K(K-1)}{2}$ because you can pair each vertex with each other vertex and divide by 2 to eliminate the same pair written backwards.

Therefore the density of a K-clique is

$$\frac{\frac{K(K-1)}{2}}{K} = \frac{K-1}{2}$$

Also, a graph that is not a K -clique has a strictly smaller density because you would have to add edges to make it a K -clique, thereby increasing the density.

Therefore if a graph has density equal to $\frac{|V|-1}{2}$ it must be a K -clique.

We construct an input to the Reasonably Sized problem as follows:

Let $D = \frac{K(K-1)}{2}$ Let G be the same graph as in the K -clique problem. Let K be the same K .

This correctly outputs YES when there is a K -clique and NO when there isn't because of the way we set D .

Since the problem is NP and NP-hard, it's NP-Complete.

3: A Well-Separated Problem

(a) The problem is in NP because there exists a poly-time verifier as follows:

Define the proof of the solution as a K -element subset of X . We compute distances between each pair of distinct points and check that they are greater than or equal to Δ . If all pairs satisfy the condition, then we verify that this is a solution. Otherwise it is not.

The Well-Separated problem is in NP-Hard because there is a Karp Reduction from the Independent Set decision problem which is NP-hard to this problem. The reduction is as follows:

Independent set

Given a graph $G = (V, E)$ and integer k ,
we want to output YES
if there exists a set of vertices of size k
such that no two of them are adjacent.

To craft our input to the Well-Separated oracle,
we construct a set X from the vertices of G ,
let $K = k$,
let $\Delta = 1.25$,
for all $i \in V$ let $d(i, i) = 0$,
for all $i, j \in V : i \neq j \wedge (i, j) \in E$ let $d(i, j) = 1$
for all $i, j \in V : i \neq j \wedge (i, j) \notin E$ let $d(i, j) = 1.5$

We pass this input to the well separated problem, which will return YES
if there exists a set of elements of size K where all distances are greater than 1.25.

Observe these elements in X map directly to vertices in V which are not adjacent due to the way we constructed the input to the Well-Separated problem.

Also note that the construction of d correctly obeys the triangle inequality, because two of the shortest distances ($1 + 1$) is still greater than the longest distance (1.5).

(b) We will present the algorithm, prove a condition about it's correctness, and show that it runs in poly time.

Algorithm:

Call one set with separation at least $\Delta^*/2$ set C .

We will maintain a vector of all points which are potentially in C , initially containing all points.

We will also maintain a vector of points which we know are in C .

1. Pick a point u potentially in C , and examine how far away all other points are from it.
2. Add u to the set of known points C .
3. Remove u from the set of points potentially in C .
4. Remove all points not at least $\Delta^*/2$ from u from the set of points potentially in C .

Repeat for a total of K iterations,
resulting in a set C with K points at least $\Delta^*/2$ away from each other
since at each step we eliminate points which could invalidate this invariant.

Now we need a proof that we can perform this K times, or in other words, we never run out of points potentially in C from which to select at each iteration.

Proof:

We were allowed to assume there exists some set of K points with separation Δ^* .
For convenience, let's call these the optimal points.

Lets examine how many optimal points are eliminated from the potentially-in- C set in one iteration of the algorithm.

Claim: At most one optimal point is removed from the potentially-in- C set.

If we select optimal point u , we will see that all the other optimal points are at least Δ^* away from u , and they will not be eliminated from the potentially-in- C set in this iteration. u will be the only optimal point removed.

If we select non-optimal point u , we will see that at most one optimal point is less than $\Delta^*/2$ away from u .

This due to the triangle inequality. Consider a nonoptimal point u , the nearest optimal point v , and any other optimal point w .

$$\Delta^* \leq d(v, w) \leq d(v, u) + d(u, w)$$

Since v no farther from u than w , $d(v, u)$ may be less than $\Delta^*/2$, but $d(u, w)$ will certainly be at least $\Delta^*/2$.

Therefore at most one optimal point is removed from the potentially-in- C set in each iteration of the algorithm.

We know there were K optimal points to start out with in the potentially-in- C set.

Therefore the algorithm can be run at least K iterations before running out of points to select.

Runtime:

The algorithm runs K iterations, each iteration takes $O(|X|)$ work, so the runtime is $O(K * |X|)$.

(c) You could run the algorithm on all possible values of Δ^* . More formally:

Modify the algorithm from B to return NONE if it runs out of points potentially in C before running K iterations.

Observe that the optimum separation delta may only be one of $\binom{|X|}{2} \leq |X|^2$ values: the distances $d(i, j) : i \neq j$.

Sort the values from highest to lowest, and run the algorithm with these values.

Return the none-NONE answer corresponding to the input with highest Δ^*

Sorting takes $O(|X|^2 * \log(|X|^2))$ and the rest takes $O(K * |X| * |X|^2)$. The limiting step is $O(K * |X|^3)$

Note: you could do this faster by using binary search instead of linear search, but it doesn't matter for the sake of this problem.