

15-451 Assignment 08

Karan Sikka

ksikka@cmu.edu

Collaborated with Dave Cummings and Sandeep Rao.

Recitation: A

November 13, 2014

1: Streaming Medians

(a)

(b)

(c)

(d)

2: Counting Substrings

3: LDIS

(a)

Choosing a random prime: Let $|\Sigma|$ be the size of the alphabet. Note that it is upper-bounded by a constant.

Turn the string into a binary string by replacing each character with a binary string of p bits where $p = \lg(|\Sigma|)$.

If the size of the original string was T , it is now t where $t = T \lg(|\Sigma|)$ which is in the order of T since $\log \sigma$ is upper bounded by a constant.

Choose a random prime between 1 and $K = 5 * p * t * \ln(pt)$.

You can do this by picking a random integer, checking if it's prime, and trying again if not. This algorithm is expected $O(\log(K))$ which is $O(\log(p) + \log(t) + \log(\log(p) + \log(t)))$ which is in $O(\log(t))$

Modified Karp-Rabin: Now we will compute the Karp-Rabin hashes of each prefix and the string of the same length following it. Starting at $i = 1$, compute $h(s[0:i])$, $h(s[i:2i])$ Increment i and compute the hashes again. Note that this takes constant time since $s[0 : i] \implies s[0 : i + 1]$ and $s[i : 2i] \implies s[i + 1 : 2i + 2]$ so to compute the hash we only have to do a constant time adjustment to the previously computed hash.

Repeat until $i > t/2$. Return the max i for which the two hash computed at the i th iteration are equal.

Proof that $\Pr[\text{false positive}] < 1/2$: Next we will show the probability of a false positive is less than $1/2$.

Let the length of the string be n . Suppose for any fixed locations i and $2i$ the probability of an incorrect match is δ . Then by a union bound over $n/2$ locations Then the probability of any incorrect match is at most $n/2 * \delta$ We want $n/2 * \delta < 1/2$ or equivalently $\delta < 1/n$.

We make an incorrect match when the hash a of one substring is equal to the hash b of another, modulo some random prime q . Formally, this occurs when $q|a - b$. $a - b$ has at most p distinct prime divisors since it's a p -bit number and each prime divisor is at least 2. If $q|a - b$, then q must have been one of the p prime divisors of $a - b$.

We want to choose a prime such that the chance that it's one of the p prime divisors is less than $1/n$. Choose K large enough so that there are pn primes between 2 and K .

If there are $\pi(x)$ primes between one and x , then $\pi(x) \geq \frac{7}{8} \frac{n}{\ln(n)}$. Thus we want to choose K such that $\pi(K) \geq \frac{7}{8} \frac{K}{\ln(K)} > pn$.

Setting $K = 5 * p * t * \ln(pt)$ achieves this result.

(b)

Construct a suffix tree for s in linear time. We will preprocess the suffix tree by running two $O(n)$ DFSs.

Do a DFS where you keep track of the length of the prefix at a node, called L , cache the index of the last occurrence of P starting on or before index $L+1$.

Ask your children, what's your last occurrence, and since your $L+1$ is strictly greater than my $L+1$, I'll filter out the invalid answers and still have the correct last occurrence before $L+1$.

```
a => yes
a a => yes
a a a => no
a a a a => no
a a a a c => no
a a a a c c => no
```

Keep running max length duplicate prefix

1. traverse a, then aa, then aaa...
2. each time, check if the following condition is true:

check if there exists another suffix prefix starting at $2 * \text{len}(\text{prefix})$ (assuming string is 1-indexed)

```
~{prefix}.*
.{len(prefix)*2}-{prefix}.*
```

3. If the condition is true, update the max-length duplicate prefix to $\text{len}(\text{prefix})$

TODO polish the above.